

# **Deep Feedforward Networks**

## **(Feedforward Neural Networks) (Multi-layer Perceptrons, MLPs)**

Yong-Sheng Chen  
Dept. Computer Science, NYCU

Sources:

“Neuroscience---exploring the brain,” 2nd ed.

by Bear, Connors, and Paradiso


Ch. 6, “Deep Learning” textbook by Goodfellow et al.

Slides by Prof. Wen-Hsiao Peng

# Human neural network vs. artificial neural network

Technology for functional brain mapping

- Functional brain mapping
- Structural analysis





- Computer vision
- CAD
- etc.

Bio-inspired deep neural network architecture


# Cortical Neurons

- 10% brain cells are neurons, 90% are glial cells
- ~ 100 billion (100 G) neurons
- ~ 2500 cm<sup>2</sup> of 2~4 mm thick sheet
- We loss >10000 neurons per day.
- 1.5 million km of fibers
- Neurons work collectively.
- 2% of adult body's weight, but 20% of its energy consumption


NVIDIA GB200:  
14,592 CUDA cores



# Neuron and Perceptron




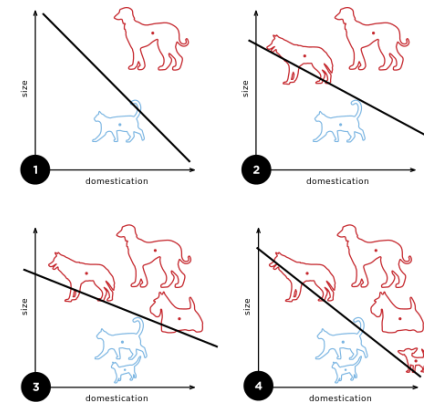
A biological neuron



An artificial neuron  
(Perceptron)  
- a linear classifier




Frank Rosenblatt  
1957



# Activation Functions


- Nonlinearity of neural network
- Binary step function

$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$




- Sigmoid function

$$f(x) = \frac{1}{1 + e^{-x}}$$

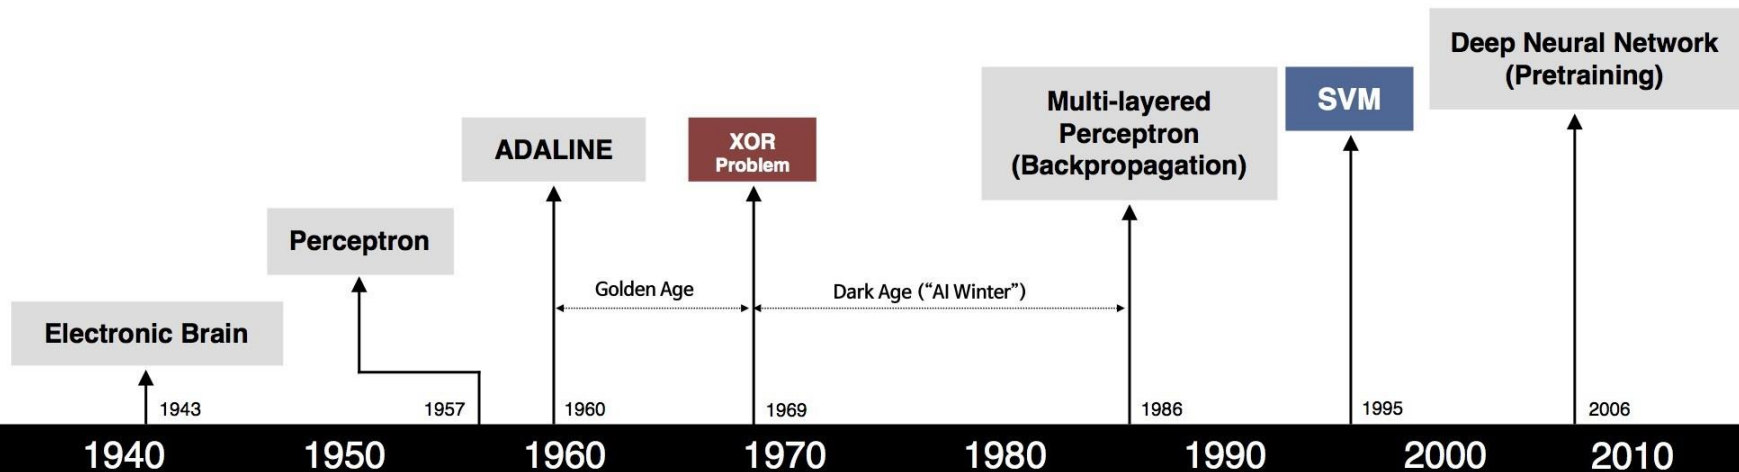


- ReLU: rectified linear unit

$$g(z) = \max\{0, z\}$$



# History of Neural Network



S. McCulloch - W. Pitts



F. Rosenblatt



B. Widrow - M. Hoff



M. Minsky - S. Papert




D. Rumelhart - G. Hinton - R. Williams




V. Vapnik - C. Cortes




G. Hinton - S. Ruslan




- Adjustable Weights
- Weights are not Learned




- Learnable Weights and Threshold




- XOR Problem



- Solution to nonlinearly separable problems
- Big computation, local optima and overfitting



- Limitations of learning prior knowledge
- Kernel function: Human Intervention




- Hierarchical feature Learning



# Multi-layer Perceptrons

- A non-linear classifier
- **Training:** find network weights  $\mathbf{w}$  to minimize the error between true training labels  $y_i$  and estimated labels  $f_{\mathbf{w}}(\mathbf{x}_i)$

$$E(\mathbf{w}) = \sum_{i=1}^N (y_i - f_{\mathbf{w}}(\mathbf{x}_i))^2$$



- Minimization can be done by gradient descent provided  $f$  is differentiable
- This training method is called **back-propagation**

D. Rumelhart, G.E. Hinton, and R. J. Williams, "Learning Internal Representations by Error Propagation," Nature, 1986.

# From linear to nonlinear

---

- To extend linear models to represent nonlinear function of  $\mathbf{x}$ :  $y = \mathbf{w}^T \mathbf{x} + b \longrightarrow y = \mathbf{w}^T \phi(\mathbf{x}) + b$ 
  - To use kernel functions such as radial basis functions (RBFs), e.g.:  $\phi(\mathbf{x}) = e^{-(\epsilon \|\mathbf{x} - \mathbf{x}_i\|)^2}$
  - Manually engineer  $\phi$ , that is, features in computer vision, speech recognition, etc.
  - To learn  $\phi$  from data:

$$y = f(\mathbf{x}; \theta, \mathbf{w}) = \mathbf{w}^T \phi(\mathbf{x}; \theta)$$




# A simple example: learning XOR

- Data:  $\chi = \left[ \begin{array}{c} 0 \\ 0 \end{array} \right], \left[ \begin{array}{c} 0 \\ 1 \end{array} \right], \left[ \begin{array}{c} 1 \\ 0 \end{array} \right], \left[ \begin{array}{c} 1 \\ 1 \end{array} \right] \}$
- Target function:  $y = f^*(\chi) = \{0, 1, 1, 0\}$
- Linear model:  $y = f(\mathbf{x}; \theta = \{\mathbf{w}, b\}) = \mathbf{x}^T \mathbf{w} + b$
- MSE loss function:



$$J(\theta) = \frac{1}{4} \sum_{\mathbf{x} \in \chi} (f^*(\mathbf{x}) - f(\mathbf{x}; \theta))^2$$

- Linear model is **NOT** able to represent XOR function.



# A simple example: learning XOR



- Use one hidden layer containing two hidden units to learn  $\phi$ .  
 $y = \mathbf{w}^T \mathbf{x} + b \longrightarrow y = \mathbf{w}^T \phi(\mathbf{x}) + b$   
 $y = f(\mathbf{x}; \mathbf{W}, \mathbf{c}, \mathbf{w}, b) = f^{(2)}(f^{(1)}(\mathbf{x}))$



# A simple example: learning XOR

- Let there be nonlinearity!
- ReLU: **R**ectified **L**inear **U**nit:  $g(z) = \max\{0, z\}$
- ReLU is applied element-wise to  $\mathbf{h}$ :

$$h_i = g(\mathbf{x}^T \mathbf{W}_{:,i} + c_i)$$



# A simple example: learning XOR

- Complete neural network model:


$$y = f(\mathbf{x}; \mathbf{W}, \mathbf{c}, \mathbf{w}, b) = f^{(2)}(f^{(1)}(\mathbf{x})) = \mathbf{w}^T \max\{0, \mathbf{W}^T \mathbf{x} + \mathbf{c}\} + b$$

- Obtain model parameters after training:

$$\mathbf{W} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}, \mathbf{c} = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, \mathbf{w} = \begin{bmatrix} 1 \\ -2 \end{bmatrix}, b = 0$$


- Run the network:

$$\mathbf{X} = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} \xrightarrow[\mathbf{W}^T \mathbf{x} + \mathbf{c}]{\mathbf{X}^T} \begin{bmatrix} 0 & -1 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{bmatrix} \xrightarrow{\text{ReLU}} \begin{bmatrix} 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{bmatrix} \xrightarrow[\mathbf{w}^T \mathbf{h} + b]{\mathbf{h}} \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$



# Gradient-Based Learning

- Nonlinearity of NN causes non-convex loss functions.
- NNs are usually trained by using iterative, gradient-based optimizers, such as **stochastic gradient descent** methods.
  - No convergence guarantee
  - Sensitive to initial values of parameters
    - Weights  $\rightarrow$  small random values
    - Bias  $\rightarrow$  zero or small positive values
  - Issues:
    - Cost (loss) functions
    - Computation of gradients
    - Gradient-based optimization



# Cost Functions: Maximum Likelihood

- Maximum likelihood model for the distribution of output  $y$ :

$$p(\mathbf{y}|\mathbf{x}; \theta)$$



Leopard cat or kitty?


- Cost function is **negative log-likelihood**, or **cross-entropy** between the training data and the model distribution:

$$J(\theta) = -E_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{y}|\mathbf{x}; \theta)$$

- If  $p_{\text{model}}(\mathbf{y}|\mathbf{x}; \theta) = N(\mathbf{y}; f(\mathbf{x}, \theta), \mathbf{I})$ ,

$$J(\theta) = \frac{1}{2} E_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{\text{data}}} \|\mathbf{y} - f(\mathbf{x}; \theta)\|^2 + \text{const}$$

- Mean squared error cost
- Large gradient and predictable for learning are preferable.
  - Many output units involve an **exp** function that can saturate when its argument is very negative.
  - log** in cost function can undoes the **exp** functions





# Output Units

- The choice of **cost function** is tightly coupled with the choice of **output unit**.
- Three common output units: **linear**, **sigmoid**, and **softmax**.
- **Linear units** for Gaussian output distributions


- Given features  $\mathbf{h}$ , output units produce a vector:

$$\hat{\mathbf{y}} = \mathbf{W}^T \mathbf{h} + \mathbf{b}$$

- Linear units are often used to produce the mean of a conditional Gaussian distribution:

$$p(\mathbf{y}|\mathbf{x}) = N(\mathbf{y}; \hat{\mathbf{y}}, \mathbf{I})$$

- Maximizing the log-likelihood is equivalent to minimizing the mean squared error.
- Linear units do not saturate.



# Output Units

---


- **Sigmoid units** for Bernoulli output distributions
  - Used to predict the value of a **binary variable**  $y$ , in classification problems with **two classes**.
  - To predict  $P(y = 1|\mathbf{x})$ , which lies in the interval  $[0,1]$ .
  - Simply use a linear unit followed by thresholding:
$$P(y = 1|\mathbf{x}) = \max\{0, \min\{1, \mathbf{w}^T \mathbf{h} + b\}\}$$
    - Not good, because the gradient would be 0 when the linear output stray outside the unit interval.

# Output Units

- **Sigmoid units** for Bernoulli output distributions

- Use a **logistic sigmoid output** combined with maximum likelihood:

$$\hat{y} = \sigma(z) = \sigma(\mathbf{w}^T \mathbf{h} + b), z = \mathbf{w}^T \mathbf{h} + b$$



- To define a probability distribution over  $y$  using  $z$ :

$$\log \tilde{P}(y) = yz, \quad y \in \{0, 1\}$$

$$\tilde{P}(y) = e^{yz}$$

$$P(y) = \frac{e^{yz}}{\sum_{y'=0}^1 e^{y'z}}$$

$$P(y) = \sigma((2y - 1)z), \quad z : \text{logit}$$

$$\begin{aligned} P(y = 0) &= \sigma(-z) \\ P(y = 1) &= \sigma(z) \end{aligned}$$

# Output Units

- **Sigmoid units** for Bernoulli output distributions

- The loss function for **maximum likelihood** learning of a Bernoulli parameterized by a **sigmoid** is:

$$J(\theta) = -\log P(y|\mathbf{x}) = -\log \sigma((2y - 1)z) = \zeta((1 - 2y)z)$$


$$\zeta(x) = \log(1 + e^x) : \text{softplus function}$$

$$y \in \{0, 1\}$$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- The softplus function does not shrink the gradient:

- It **saturates only when the answer is right** ( $y=1$ ,  $z$  is positive or  $y=0$ ,  $z$  is negative), such that  $(1-2y)z$  is negative.
- When the answer is wrong, the softplus function returns  $\sim |z|$



# Output Units

---


- Sigmoid units for Bernoulli output distributions
  - When sigmoid units are combined with other cost functions like mean squared error, the gradient vanishing problem may occur.
  - The choice of output units is tightly coupled with that of cost functions.

# Output Units

- **Softmax units** for Multinoulli output distributions
  - A probability distribution over a discrete variable with  $n$  possible values
  - Used as the output of a **classifier for  $n$  classes**
  - Goal: a vector  $\hat{\mathbf{y}}$ ,  $\hat{y}_i = P(y = i|\mathbf{x})$ 
    - Each  $\hat{y}_i$  lies in the interval  $[0,1]$ .
    - Entire vector  $\hat{\mathbf{y}}$  sums to 1.

$$\mathbf{z} = \mathbf{W}^T \mathbf{h} + \mathbf{b}, \quad z_i = \log \tilde{P}(y = i|\mathbf{x})$$

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

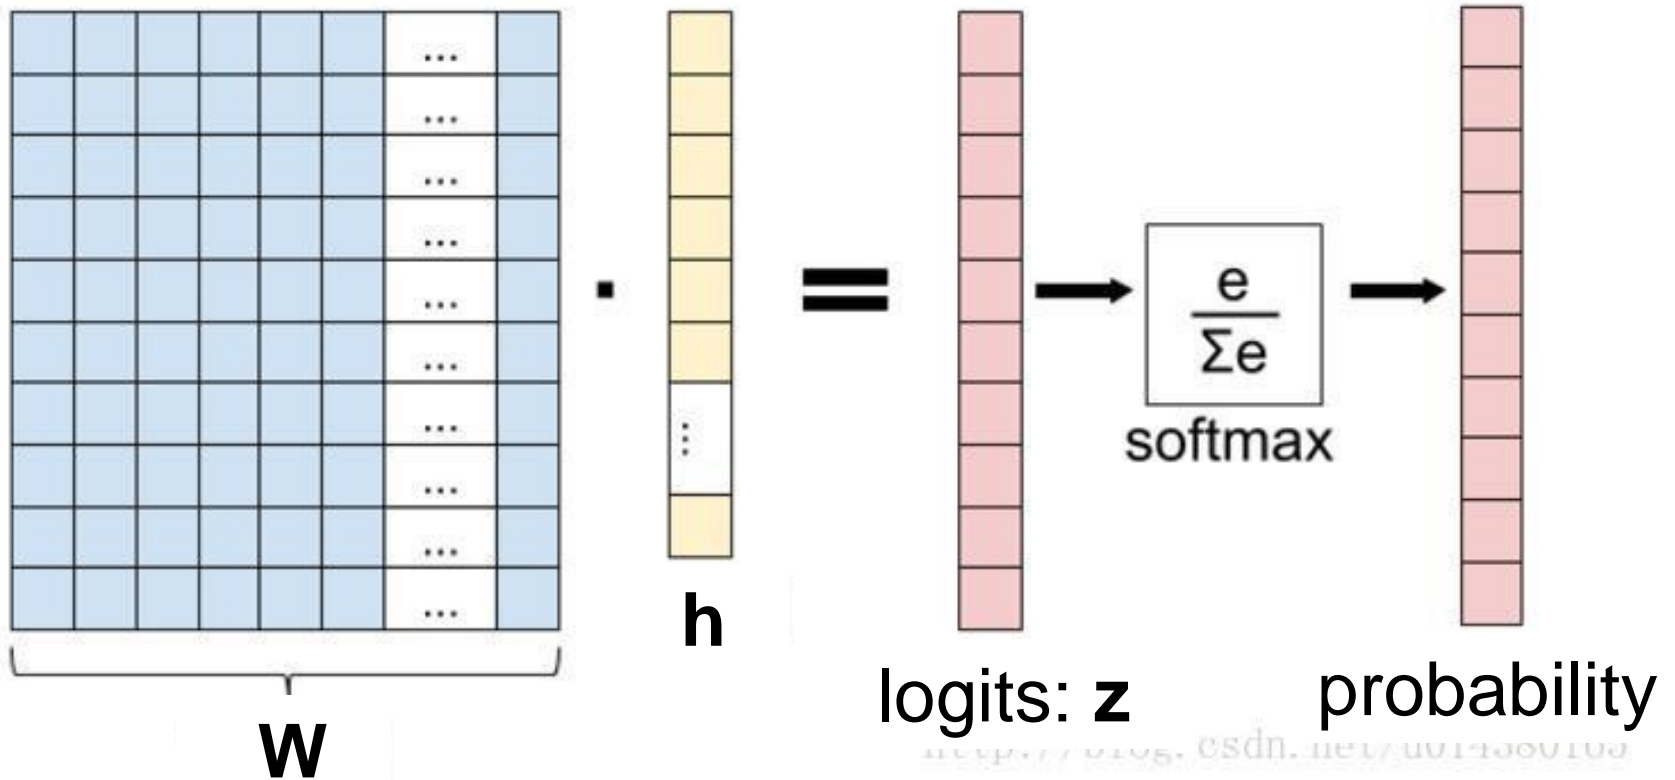


In Sigmoid unit:  $P(y = 1) = \frac{e^z}{\sum_{y'=0}^1 e^{y'z}}$



# Output Units

- Softmax units for Multinoulli output distributions



# Output Units

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

- **Softmax units** for Multinoulli output distributions

- Minimize negative log-likelihood:

$$-\log P(y = i; \mathbf{z}) = -\log \text{softmax}(z)_i = -(z_i - \log \sum_{j=1}^n e^{z_j})$$

- Input  $z_i$  always has a direct contribution to the cost function and cannot saturate.
- While minimizing, the first term encourages  $z_i$  to be pushed up, while the second term encourages all of  $\mathbf{z}$  to be pushed down.

- For the correctly classified sample,
$$z_i \gg z_k, k \neq i \rightarrow z_i - \log \sum_{j=1}^n e^{z_j} \approx z_i - z_i = 0$$

- Negative log-likelihood cost function always **strongly penalizes the most active incorrect prediction**.
- Most objective function other than the negative log-likelihood do not work well with the softmax function.

# Output Units

- **Softmax units** for Multinoulli output distributions

- Softmax function can be easily out of range for extreme input values.

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

- Note that  $\text{softmax}(\mathbf{z}) = \text{softmax}(\mathbf{z} + c)$
- We can derive a numerically stable variant of softmax:

$$\text{softmax}(\mathbf{z}) = \text{softmax}(\mathbf{z} - \max_j z_j)$$

- An output  $\text{softmax}(\mathbf{z})_i$  saturates to 1 when the corresponding input  $z_i$  is maximal and much greater than all other inputs.
- An output  $\text{softmax}(\mathbf{z})_i$  can also saturate to 0 when  $z_i$  is not maximal and the maximum is much greater.

# Output Units

- **Softmax units** for Multinoulli output distributions
  - Softmax is a way to create a form of competition between the units that participate in it.
  - From a neuroscience point of view, lateral inhibition is believed to exist between nearby neurons, that is, winner-take-all.
  - Softmax
    - softened version of arg max
    - Continuous and differentiable

*Nature Neuroscience* 2, 375 - 381 (1999)  
doi:10.1038/7286


Attention activates winner-take-all competition among visual filters

D.K. Lee<sup>1,2</sup>, L. Itti<sup>1,2</sup>, C. Koch<sup>1</sup> & J. Braun<sup>1</sup>


Shifting attention away from a visual stimulus reduces, but does not abolish, visual discrimination performance. This residual vision with 'poor' attention can be compared to normal vision with 'full' attention to reveal how attention alters visual perception. We report large differences between residual and normal visual thresholds for discriminating the orientation or spatial frequency of simple patterns, and smaller differences for discriminating contrast. A computational model, in which attention activates a winner-take-all competition among overlapping visual filters, quantitatively accounts for all observations. Our model predicts that the effects of attention on visual cortical neurons include increased contrast gain as well as sharper tuning to orientation and spatial frequency.

# Hidden Units

- Design of hidden units in an extremely active area of research.
- ReLUs are an excellent default choice.
- Although not differentiable at all point, it is still okay to use for gradient-based learning algorithm.
  - Use left or right derivative, instead.
- Hidden units compute:
  - An affine transformation  $z = \mathbf{W}^T \mathbf{x} + \mathbf{b}$
  - An element-wise nonlinear function  $g(\mathbf{z})$



# ReLU




- **ReLU**: rectified linear unit:  $g(z) = \max\{0, z\}$
- Easy to optimize because the derivatives through ReLU remain large whenever the unit is active.
- When initializing the parameters of the affine transformation, it can be a good practice to set all elements of  $\mathbf{b}$  to a small, positive value, such as 0.1. (why?)  $\mathbf{h} = g(\mathbf{W}^T \mathbf{x} + \mathbf{b})$
- Drawback: ReLUs cannot learn via gradient-based methods on examples with zero activation.



# Generalizations of ReLUs

- Use a non-zero slope  $\alpha_i$  when  $z_i < 0$ :  
$$h_i = g(\mathbf{z}, \alpha)_i = \max\{0, z_i\} + \alpha_i \min\{0, z_i\}$$
- Absolute value rectification fixes  $\alpha_i = -1$ :  $g(z) = |z|$
- **Leaky ReLU** fixes  $\alpha_i$  to a small value like 0.01.
- **Parametric ReLU** or **PReLU** treats  $\alpha_i$  as a learnable parameter.




# Generalizations of ReLUs

- Maxout units:


$$g(\mathbf{z})_i = \max_{j \in [1, k]} z_{ij} = \max\{\mathbf{w}_1^T \mathbf{x} + \mathbf{b}_1, \mathbf{w}_2^T \mathbf{x} + \mathbf{b}_2, \dots, \mathbf{w}_k^T \mathbf{x} + \mathbf{b}_k\}$$

- It becomes an ReLU when  $k=2$ ,  $\mathbf{w}_1=\mathbf{b}_1=0$
- It can **learn** a piecewise linear, convex activation function with up to  $k$  pieces.



# Maxout Units

- Maxout unit learns activation function
- Maxout unit improves abstraction ability





# Maxout Inspired by Lateral Inhibition

- The **pathway** along which the signal flows that determines the *functionality* of information processing


(Kandel et al. Principles of neural science)

- Lateral inhibition



# Architecture Design

- Architecture: overall structure of the network
  - How many units it should have
  - How these units should be connected to each other
  - How to choose the depth and width of each layer
- Deeper networks often:
  - Use far fewer units per layer and far fewer parameters
  - Generalize to the test set
  - Are harder to optimize



# Universal Approximation Properties and Depth


---

- The **universal approximation theorem** states that a feedforward network with one linear output layer and **at least one hidden layer** with any “squashing” activation function (e.g., sigmoid) can approximate any Borel measurable function, provided that the network is given enough hidden units.
- Any continuous function  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$  on a closed and bounded subset of  $\mathbb{R}^n$  is Borel measurable.
- A feedforward network with a single layer is sufficient to represent any function, but the layer may be infeasibly large and may fail to learn and generalize correctly.
- Using deeper models can reduce the number of units required to represent the desired function and can reduce the generalization error.




# Universal Approximation Properties and Depth

- A network with absolute value rectification creates mirror images of the function computed on top of hidden units.




## Better Generalization with Greater Depth

- Empirically, deeper networks generalize better.
- SVHN

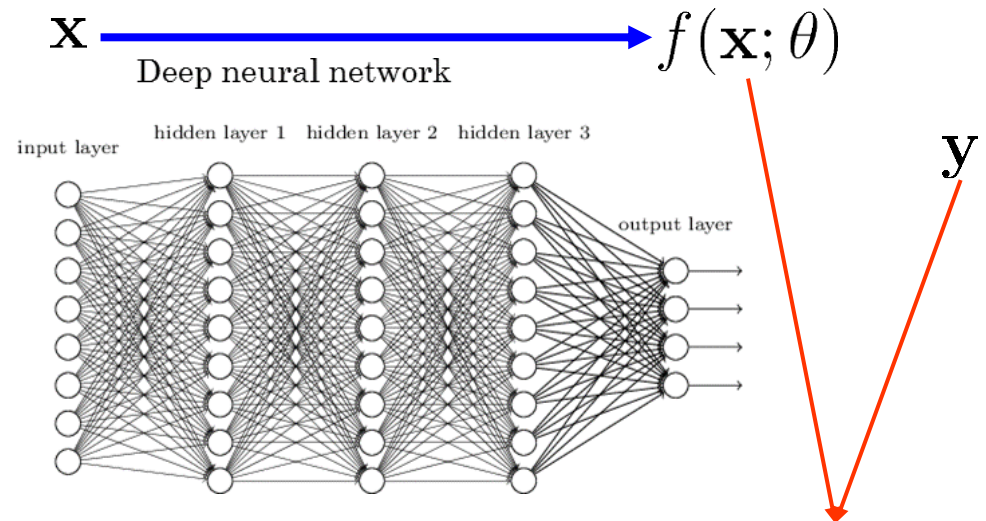



## Large, Shallow Models Overfit More




# Back-Propagation

- During inference:  
(forward propagation)



- During training:
- Backpropagation:




- Stochastic gradient descent is used to perform the learning using these gradients.

# Computational Graphs



- Each node indicate a variable

- Scalar
- Vector
- Matrix
- Tensor


(multi-dimensional array)



(a)

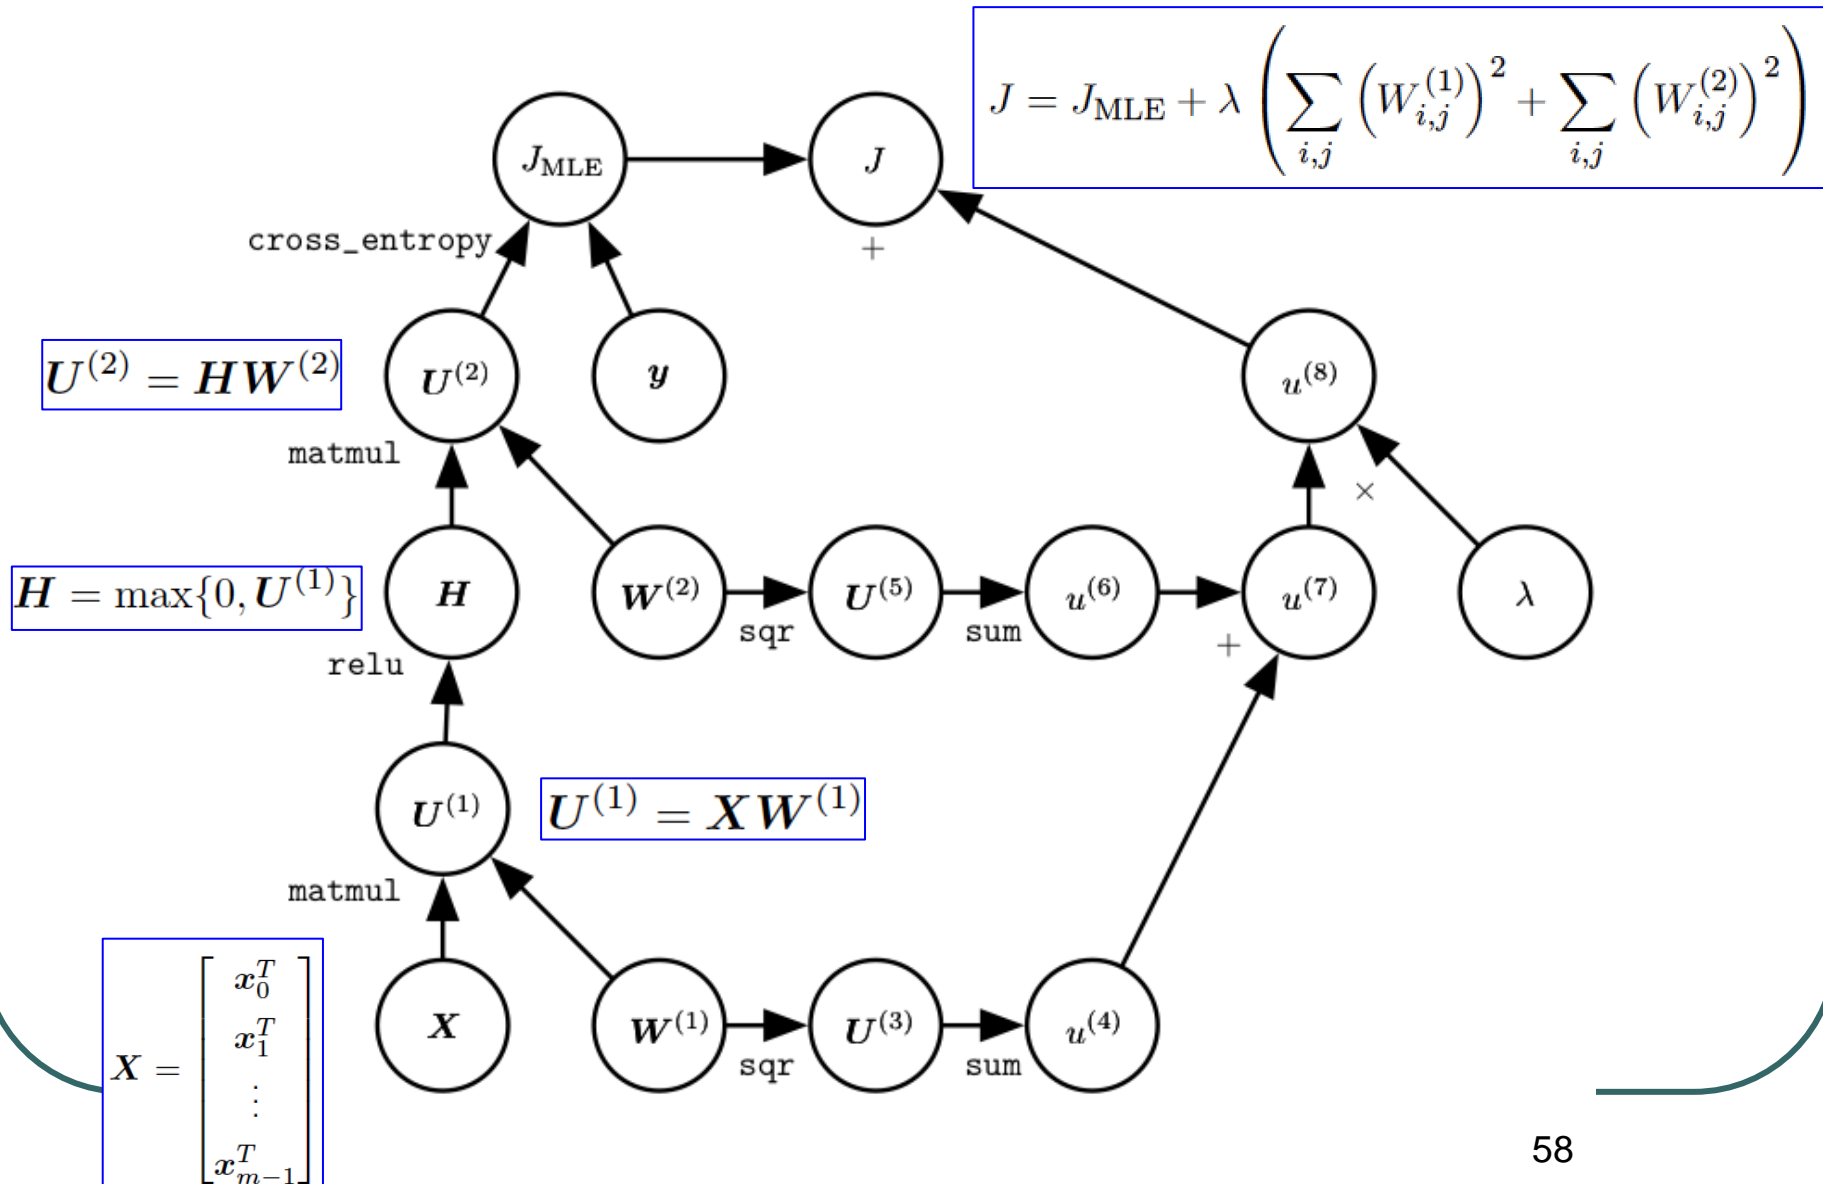


$$\mathbf{H} = \max\{0, \mathbf{W}^T \mathbf{X} + \mathbf{b}\}$$



(d)

# Backprop for MLP Training



- Backprop: To compute  $\nabla_{\mathbf{W}^{(1)}} J$  and  $\nabla_{\mathbf{W}^{(2)}} J$ 
  - Two paths from  $J$  to the weights (only one path illustrated)
  - Assume  $\nabla_{\mathbf{U}^{(2)}} J = \mathbf{G}$
  - Then  $\nabla_{\mathbf{W}^{(2)}} J = \mathbf{H}^T \mathbf{G}$  (cf. **Example 1** in Matrix Case)
  - Similarly,  $\nabla_{\mathbf{H}} J = \mathbf{G} \mathbf{W}^{(2)T}$  (cf. **Example 2** in Matrix Case)
  - Tracing back further, we have  $\nabla_{\mathbf{U}^{(1)}} J = \mathbf{G}'$  by zeroing out elements in  $\nabla_{\mathbf{H}} J$  corresponding to entries of  $\mathbf{U}^{(1)}$  less than zero
  - Again,  $\nabla_{\mathbf{W}^{(1)}} J = \mathbf{X}^T \mathbf{G}'$  (cf. **Example 1** in Matrix Case)

### Example 1

$$Y = g(X) = WX$$

$$\nabla_X z = W^T \nabla_Y z$$

### Example 2

$$Y = g(X) = XW$$

$$\nabla_X z = (\nabla_Y z) W^T$$

$$U^{(2)} = HW^{(2)}$$

$$U^{(1)} = XW^{(1)}$$

